



# Network Programming with Boost

Jeff Garland

[jeff@crystalclearsoftware.com](mailto:jeff@crystalclearsoftware.com)

# Tutorial Goals

- ◆ Describe the basics of writing networked programs
- ◆ Writing network enabled programs including:
  - understanding of some of the issues in network programming
  - an understanding of 'asynchronous programming'
  - basics of TCP based sockets
  - integration of sockets and i/o streams
  - understanding of using timers in network programming
- ◆ Using serialization with asio to build a general object messaging frameworks

# Topics Covered

- ◆ Network programming
  - Terminology and theory
  - Tools
- ◆ Boost Libraries
  - asio
  - serialization
  - bind
  - regex
  - smart\_ptr
  - date-time

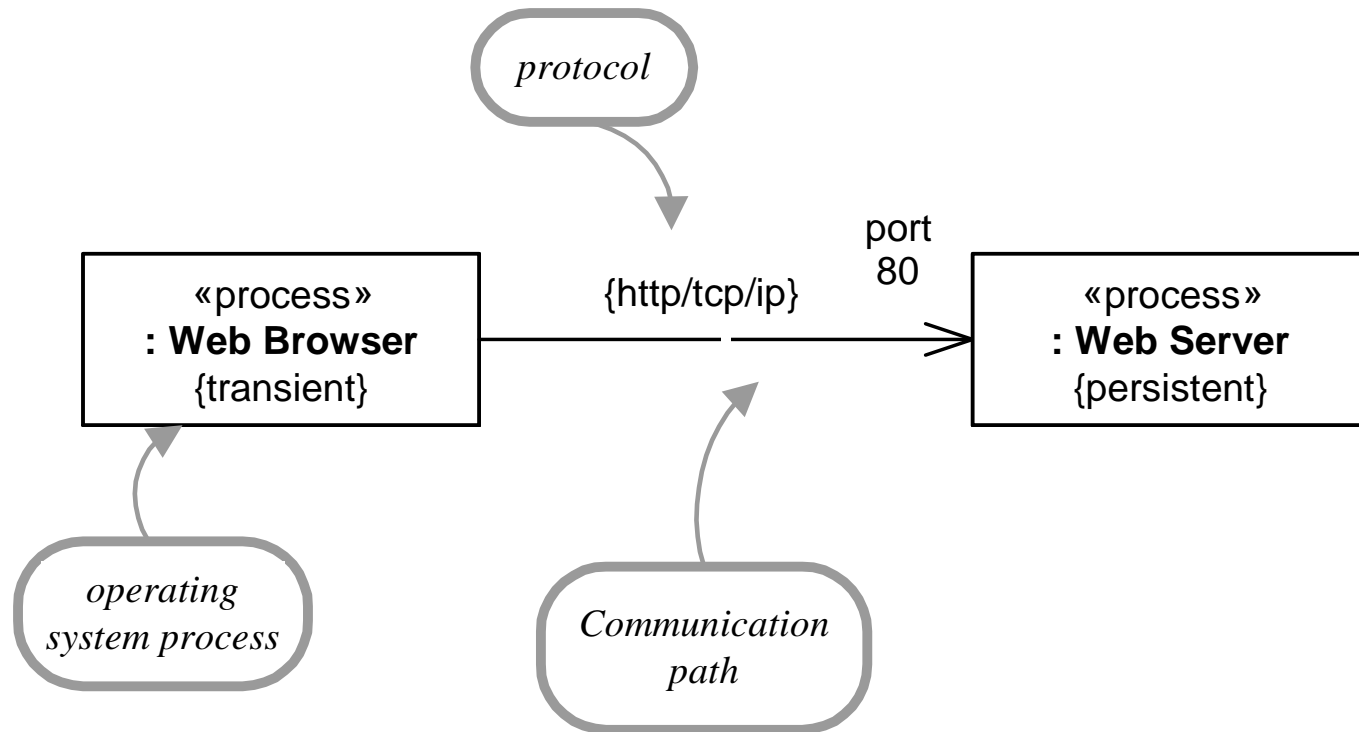
# Session Agenda

- ◆ Introduction
- ◆ Boost.asio Overview
- ◆ Basic Synchronous Client Program
- ◆ Boost.asio in depth
- ◆ Asynchronous Server Program
- ◆ Boost.serialization – short overview
- ◆ Pitfalls of Network Programming
- ◆ Network Programming Tools
- ◆ Resources and Conclusion

# What's the need?

- ◆ Many modern applications are now distributed applications
- ◆ Networking Advancements
  - Ubiquitous networks - opens possibility of making more apps networked
  - 'Fast' networks
    - Gigabit ethernet
    - Wireless N
- ◆ Protocol Development
  - Continued advancements in 'application protocols' (example RSS)

# Basic Client Server Communication

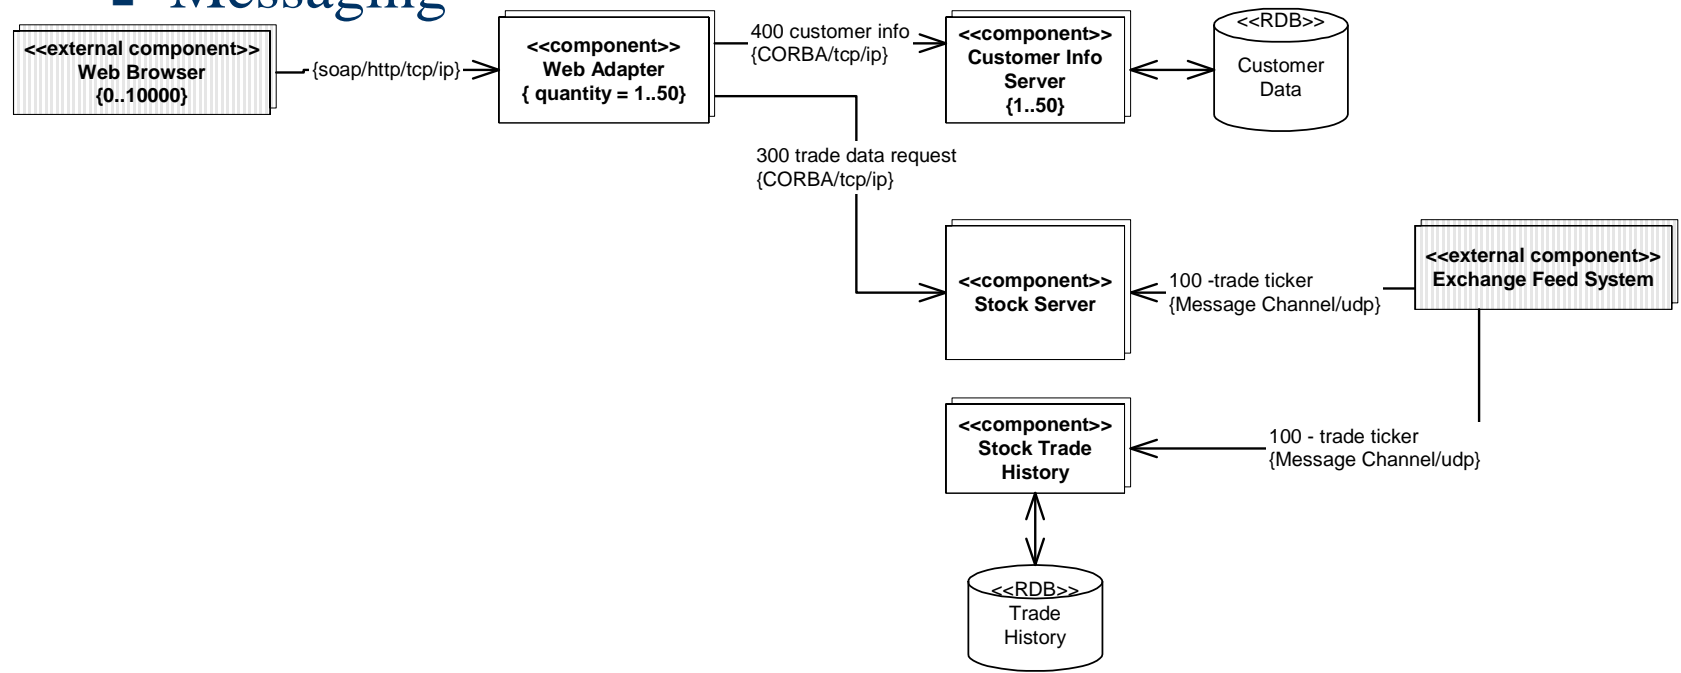


# Networking Terminology

- ◆ **Client** – Process that accesses a remote service.
- ◆ **Server** – Process that provides a service for remote computers.
- ◆ **Runtime Component** - a construct that groups runtime entities and provides/consumes a set of interfaces.
- ◆ **Process** - an instance of an operating system process

# Types of Network Communication - Example

- ◆ Web page to display current price of stock
  - Client – server type
- ◆ Exchange feeds
  - Messaging





# Types of Network Communications

- ◆ Service
  - Client/Server
  - RPC
  - Client connects to server, makes a request, optionally receives response
  - Suitable for larger data transfers
  - Point to Point
- ◆ Messaging
  - Publish subscribe, 'Multi-cast'

# Networking Concepts

- ◆ Connectionless Protocol
  - UDP/IP
    - UDP – User Datagram Protocol
    - IP – Internet Protocol
- ◆ Connection-oriented Protocol
  - TCP/IP
    - TCP - Transmission Control Protocol
  - SCTP/IP
    - SCTP – Stream Control Transmission Protocol

# IP – Internet Protocol

- ◆ Base addressing scheme
  - Used for almost all communications now
- ◆ IPv4
  - Typical address: 192.168.1.100
  - 32 bits
- ◆ IPv6
  - Typical address:  
3ffe:0501:0008:0000:0260:97ff:fe40:efab
  - 128 bits
- ◆ Special addresses
  - Loopback 127.0.0.1 – localhost
  - 192.xxx.xxx.xxx

# Session Agenda

- ◆ Introduction
- ◆ Boost.asio Overview
- ◆ Basic Synchronous Client Program
- ◆ Boost.asio in depth
- ◆ Asynchronous Server Program
- ◆ Boost.serialization – short overview
- ◆ Pitfalls of Network Programming
- ◆ Network Programming Tools
- ◆ Resources and Conclusion

# asio 101 – What is asio?

- ◆ Boost cross-platform system OS access library
  - Networking is just one aspect
  - Timers
  - May eventually support file access and others
- ◆ What does asio mean?
  - Asynchronous Input-Output...nope
  - Actually - play on words Australian Security Intelligence Organisation
- ◆ Pronounced - ay-zee-oh
- ◆ Library written by Christopher M. Kohlhoff

# asio 101 – history / info

- ◆ Started outside of Boost
  - Was quite successful as a standalone project
- ◆ Reviewed and accepted into Boost Dec/Jan 2005
- ◆ In Boost 1.35 – just released

# asio 101 – Design Goals

- ◆ Efficient - Allow for minimal data copying
- ◆ Scalable – Up to thousands of concurrent connections
- ◆ Integrate well with standard library
  - streambuf integration
  - Foundation for TR2 Networking Proposal
    - <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2006/n2054.pdf>
- ◆ Easy to learn
  - Network concepts based on Berkley sockets
- ◆ Portable

# asio 101 - Platforms

- ◆ Win32 with VC++ 7.1 and 8.x
- ◆ Win32 with Borland C++ builder
- ◆ Win32 with MinGW or Cygwin / g++
- ◆ Linux with g++ 3.3 or later
- ◆ Solaris with g++3.3 or later
- ◆ Mac OS X 10.4 using g++ 3.3 or later
- ◆ FreeBSD with g++3.3 or later
- ◆ QNX Neutrino 6.3 with g++ 3.3. or later



# asio 101 – lib dependencies

- ◆ Other boost library dependencies
  - ‘system’ used for integrated error handling
  - date\_time used for time specifications
    - Technically optional

# asio 101 - header Only?


- ◆ Review version of asio was header-only
- ◆ Newer versions depend on boost.system
  - Compiled library
- ◆ List discussion of merits

# Session Agenda

- ◆ Introduction
- ◆ Boost.asio Overview
- ◆ Basic Synchronous Client Program
- ◆ Boost.asio in depth
- ◆ Asynchronous Server Program
- ◆ Boost.serialization – short overview
- ◆ Pitfalls of Network Programming
- ◆ Network Programming Tools
- ◆ Resources and Conclusion

# Fetch a web page

Service Address  
209.131.36.158:80



# Example: fetch a web page

- ◆ Simple form of ‘wget’
- ◆ Steps:
  - Step 1: Includes and setup stuff
  - Step 2: Get the address of the server
  - Step 3: Make a connection to server
  - Step 4: Write http request to server
  - Step 5: Read resulting data
    - 5a – read the header return
    - 5b – process the data

# Functions needed in Networked Apps

- ◆ Find address of service
- ◆ If connection oriented:
  - Connect to service, or wait for connections
  - Reconnect on failures
- ◆ Deal with Data
  - Format request or response data to transmit on wire
  - Protocol support
  - Buffering
- ◆ Handle Failures
- ◆ Multiplexing and Timing
- ◆ Setup connection and other options

# Fetch Web Page 1: Setup

```
#include <boost/asio.hpp>
#include <boost/regex.hpp> //not needed in latest
#include <boost/system/system_error.hpp> //latest
    only
#include <string>
#include <iostream>
using namespace boost::asio;
using namespace boost::asio::ip;
using std::cout;
using std::endl;
using std::string;
```

# Fetch Web Page 1:Setup - Main

```
int main (int argc, char* argv[]) {  
    try  
    {  
        boost::asio::io_service io_service;  
        //To Be expanded  
    }  
    catch (std::exception& e)  
    {  
        std::cout << "Exception: " << e.what() << std::endl;  
    }  
  
    return 0;  
}
```



# Fetch Web Page 2: Get Address

```
boost::asio::io_service io_service;
```

```
//eg: www.boost.org
```

```
string host_name = "localhost";
```

```
tcp::resolver resolver(io_service);
```

```
tcp::resolver::query query(host_name, "http");
```

```
tcp::resolver::iterator endpoint_iterator = resolver.resolve(query);
```

```
tcp::resolver::iterator end;
```

```
// To be expanded
```

```
while (endpoint_iterator != end)
```

# Fetch Web Page 3: Connect Socket

```
tcp::socket socket(io_service);

//connect error is exception type!
//boost::asio::error connect_error = boost::asio::error::host_not_found;
boost::system::error_code connect_error = boost::asio::error::host_not_found;
while (connect_error && endpoint_iterator != end)
{
    socket.close();
    socket.connect(*endpoint_iterator++, connect_error);
}

if (connect_error)
    throw connect_error;
//To Be expanded
```

# Fetch Web Page 4: Write Fetch Request

```
string page_path = "/index.html";

boost::asio::streambuf request;
std::ostream request_stream(&request);
//http protocol stuff...
request_stream << "GET " << page_path << " HTTP/1.0\r\n";
request_stream << "Host: " << host_name << "\r\n";
request_stream << "Accept: */*\r\n";
request_stream << "Connection: close\r\n\r\n";

// Send the request.
boost::asio::write(socket, request);

//To Be expanded
```

# Fetch Web Page 5: Receive Data – Check Header

Read the response status line.

```
boost::asio::streambuf response;  
boost::asio::read_until(socket, response, "\r\n");  
//boost::asio::read_until(socket, response, boost::regex("\r\n")); //v 0.3.7
```

// Check that response is OK.

```
std::istream response_stream(&response);  
std::string http_version;  
response_stream >> http_version;  
unsigned int status_code;  
response_stream >> status_code;  
std::string status_message;  
std::getline(response_stream, status_message);
```

# Fetch Web Page 5: Receive Data – Check Header

```
//....continued...
if (!response_stream || http_version.substr(0, 5) != "HTTP/")
{
    std::cout << "Invalid response\n";
    return 1;
}
if (status_code != 200)
{
    std::cout << "Response returned with status code " << status_code << "\n";
    return 1;
}

//To Be expanded
```

# Fetch Web Page 6: Receive Data

```
// Read the response headers, which are terminated by a blank line.  
//boost::asio::read_until(socket, response, boost::regex("\r\n\r\n")); //v 0.3.7  
boost::asio::read_until(socket, response, "\r\n\r\n");
```

```
// Process the response headers.  
string header;  
cout << "**** Header ****" << endl;  
while (std::getline(response_stream, header) && header != "\r")  
{  
    cout << header << "\n";  
}  
cout << "\n";
```

# Fetch web Page 6: Receive Data

```
cout << "**** Part 1 ****" << endl;
// Write whatever content we already have to output.
if (response.size() > 0)
    cout << &response;

cout << "**** Rest ****" << endl;
boost::system::error read_error;
// Read until EOF, writing data to output as we go.
while (boost::asio::read(socket,
                        response,
                        boost::asio::transfer_at_least(1),
                        read_error)) {
    // boost::asio::assign_error(read_error))) {
    cout << &response;
}
if (read_error != boost::asio::error::eof)
    throw read_error;
```

# Session Agenda

- ◆ Introduction
- ◆ Boost.asio Overview
- ◆ Basic Synchronous Client Program
- ◆ Boost.asio in depth
- ◆ Asynchronous Server Program
- ◆ Boost.serialization – short overview
- ◆ Pitfalls of Network Programming
- ◆ Network Programming Tools
- ◆ Resources and Conclusion



# asio in depth - namespace structure

- ◆ Base Namespace: `boost::asio`
  - Core concept templates
    - `basic_socket`, `basic_deadline_timer`, etc
  - Contains `io_service`, `buffer`, others
  - Various free functions (read/write)
- ◆ `boost::asio::ip` - concrete ip related implementations
  - Defines core protocol level classes
    - `tcp` – for tcp related capabilities
      - ◆ `acceptor`, `endpoint`, `iostream`, `resolver`, `socket`
    - `udp` – for udp related capabilities
      - ◆ `endpoint`, `socket`, `resolver`
  - Sub namespace for multicast and unicast options
  - Some concept templates as well -- `basic_endpoint` (moved since v0.3.7)
- ◆ `boost::asio::ssl` – Secure sockets
  - `context`, `context_base`, `stream_base` (etc)
- ◆ `boost::asio::placeholders` – bind placeholders
  - `error`, `bytes_transferred`, `iterator`

# asio in depth – sync read

## ◆ Synchronous read functions

- Attempt to read a certain amount of data from a stream before returning.
- Many variants:
  - read
  - read\_until

```
template<typename Sync_Read_Stream,  
         typename Mutable_Buffers>
```

```
std::size_t boost::asio::read (Sync_Read_Stream &s,  
                               const Mutable_Buffers &buffers)
```

## asio in depth – sync write

- ◆ Write all of the supplied data to a stream before returning.
  - ◆ Several variants with different buffer types
- ```
template<typename Sync_Write_Stream,  
        typename Const_Buffers>  
std::size_t boost::asio::write (Sync_Write_Stream &s,  
                                const Const_Buffers &buffers)
```

# Asio : ip address

- ◆ Class that represents a v4 or v6 ip address
- ◆ Example: classic v4 address:
  - 192.168.100.1
- ◆ Couple variants if you know the ip version
  - address\_v4
  - address\_v6

# Key ip\_address functions

| Method                                | Parameters  | Return      | Description                      |
|---------------------------------------|-------------|-------------|----------------------------------|
| address                               |             |             | Default constructor.             |
| to_string                             |             | std::string | Convert to string.               |
| from_string                           | const char* | address     | Construct address from a string. |
| operator==<br>operator!=<br>operator< | Address     | bool        | Less than comparable functions   |

- ◆ Several other variants of from\_string
- ◆ To\_v4 and to\_v6 if you need specific address type

# Other functions: Get our Hostname

```
#include <boost/asio.hpp>
#include <iostream>
using namespace boost::asio;
int main()
{
    try
    {
        std::string host_name = ip::host_name();
        std::cout << "hostname: " << host_name << std::endl;
    }
    catch (std::exception&)
    {}
    return 0;
}
```

# asio concepts: Endpoint

- ◆ Combines protocol and address of a service
- ◆ Use by asio in variety of functions:
  - Server - listening for clients
  - Client - connecting to a server

# Key endpoint functions

| Method                  | Description                                                                                    |
|-------------------------|------------------------------------------------------------------------------------------------|
| Constructor             | <code>basic_endpoint(const Protocol &amp;protocol, unsigned short port_num)</code>             |
| Constructor for connect | <code>basic_endpoint(const boost::asio::ip::address &amp;addr, unsigned short port_num)</code> |
|                         |                                                                                                |



# asio concepts: Resolver

- ◆ Resolver
  - Provides the ability to retrieve an endpoint(s) from name
  - Convert yahoo.com to an ip address
- ◆ Provides endpoint iterator

# Key resolver functions

| Method        | Description                                                                                          |
|---------------|------------------------------------------------------------------------------------------------------|
| Constructor   | <code>basic_resolver(io_service&amp;)</code>                                                         |
| Resolve       | <code>iterator resolve (const endpoint_type &amp;e)</code> Resolve an endpoint to a list of entries. |
| Resolve query | <code>iterator resolve( const query &amp; q, boost::system::error_code &amp; ec)</code>              |

# asio concepts: basic\_socket

- ◆ Provides socket capabilities
- ◆ Concrete types for tcp and udp
  - `boost::asio::ip::tcp::socket`
  - `boost::asio::ip::udp::socket`

# Socket – Getting - Setting Options

```
//set nodelay option
```

```
boost::asio::tcp::socket socket (io_service);
```

```
boost::asio::tcp::no_delay option(true);
```

```
socket.set_option(option);
```

```
//get no_delay setting
```

```
boost::asio::tcp::no_delay option;
```

```
socket.get_option(option);
```

```
bool is_set = option.get();
```

# Key Socket Functions -

| Method        | Description                                                    |
|---------------|----------------------------------------------------------------|
| async_connect | Start a connect to an endpoint                                 |
| connect       | Start a connection - synchronous                               |
| cancel        | Cancel all asynchronous operations associated with the socket. |
| close         | Close the socket.                                              |
| open          | Open the socket using the specified protocol.                  |
| native        | Get the underlying socket representation (non portable)        |

# Key Socket Functions - Info

| Method          | Description                                                     |
|-----------------|-----------------------------------------------------------------|
| remote_endpoint | Get the other ip information from the other side of the socket. |
| local_endpoint  | IP information from local side.                                 |

# asio concepts: io\_service

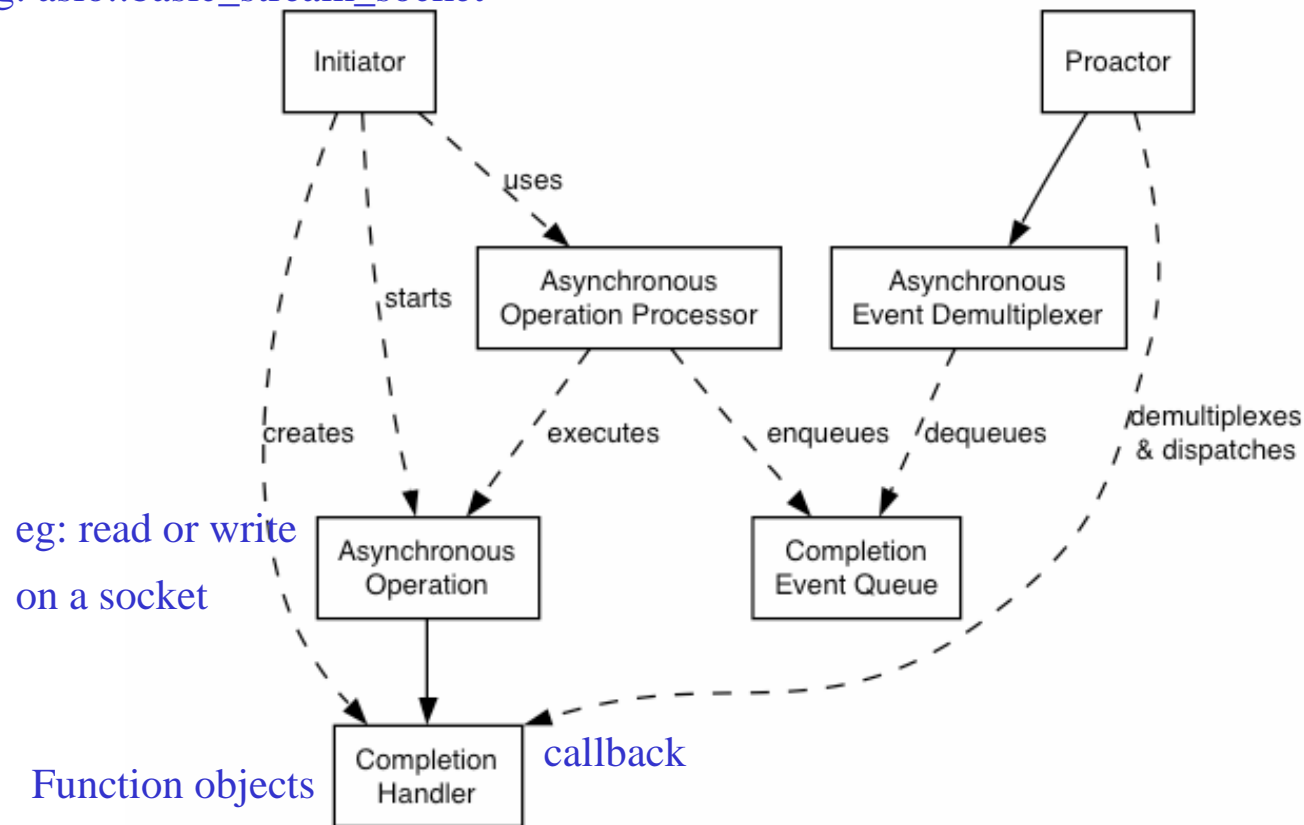
- ◆ Provides an ‘event loop’
  - Callback functions are registered
  - Requests initiated
  - Callbacks de-multiplexed and dispatched to handlers
- ◆ Allows for multi threaded operations
  - Calling run from multiple threads creates pool
  - Handler registration and thread NOT associated

# Event De-multiplexing

Your application

Eg: asio::basic\_stream\_socket

boost::asio::io\_service





# Key io\_service functions

| Method               | Description                                                                                                                    |
|----------------------|--------------------------------------------------------------------------------------------------------------------------------|
| io_service           | Default constructor.                                                                                                           |
| run                  | Start event processing - blocks until all work has finished and there are no more handlers to be dispatched or stop is called. |
| Stop (was interrupt) | Stop the even loop. Returns immediately, there may be a delay before run returns.                                              |

# asio concepts: timing – deadline\_timer

- ◆ Timers a frequent need in network programming
  - Example: retry connection on failure
- ◆ Deadline timer provides this capability
  - Sync and async versions

# Deadline timer sync

```
using boost::posix_time::seconds;
using boost::posix_time::milliseconds;

// Construct a timer without setting an expiration time.
boost::asio::deadline_timer timer(io_service);

// Set an expiration time relative to now.
timer.expires_from_now(seconds(5));
// timer.expires_from_now(milliseconds(50));

// Wait for the timer to expire.
timer.wait();
```

# Deadline Timer - async

```
void handler(const boost::asio::error& error) {
    if (!error)
    {
        // The current time is:
        boost::posix_time::ptime now(microsec_clock::local_time());
        std::cout << "Timeout callback: " << now << std::endl;
        // Timer expired.
    }
}

// Construct a timer with an absolute expiration time.
boost::asio::deadline_timer timer(io_service,
                                   boost::posix_time::time_from_string("2005-12-07 23:59:59.000"));
// boost::asio::deadline_timer timer(io_service, seconds(3)); //alternative for relative time

// Start an asynchronous wait.
timer.async_wait(handler);
```

# Session Agenda

- ◆ Introduction
- ◆ Boost.asio Overview
- ◆ Basic Synchronous Client Program
- ◆ Boost.asio in depth
- ◆ Asynchronous Server Program
- ◆ Boost.serialization – short overview
- ◆ Pitfalls of Network Programming
- ◆ Network Programming Tools
- ◆ Resources and Conclusion

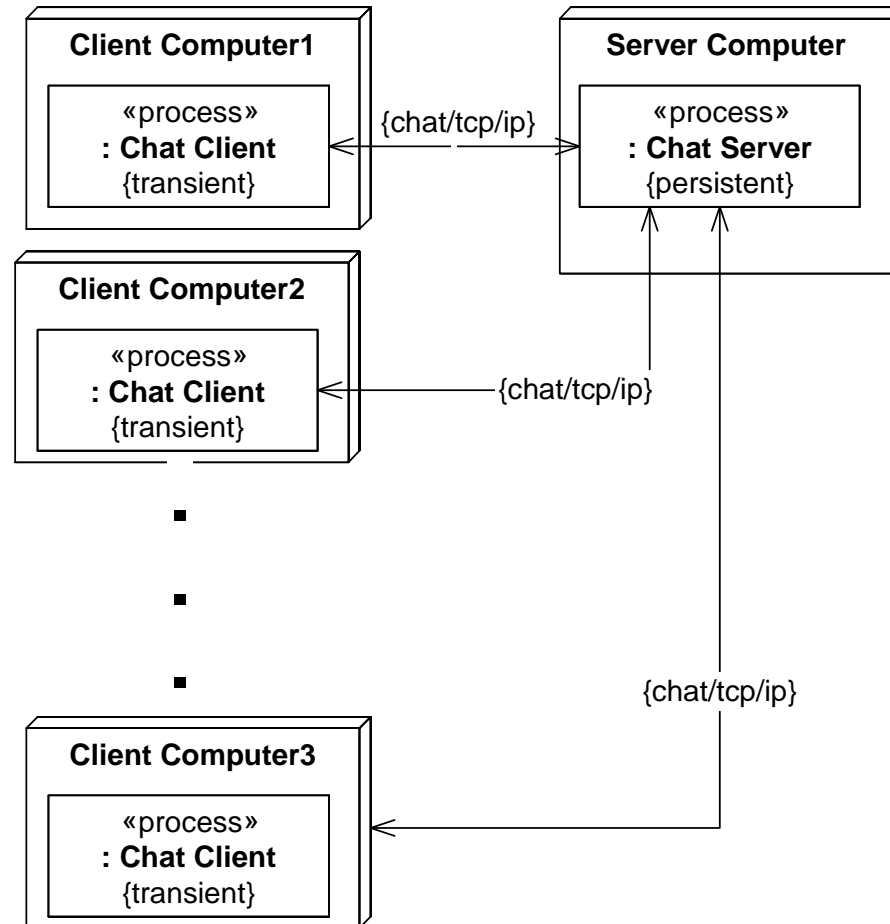
# Building Chat Server

- ◆ Tasks include
  - Listen for a client to connect
  - ‘accept’ connection requests
  - Handle data from connections
    - When received send messages back to all clients
- ◆ Use async functions
- ◆ Use Chat program

# Acceptor pattern

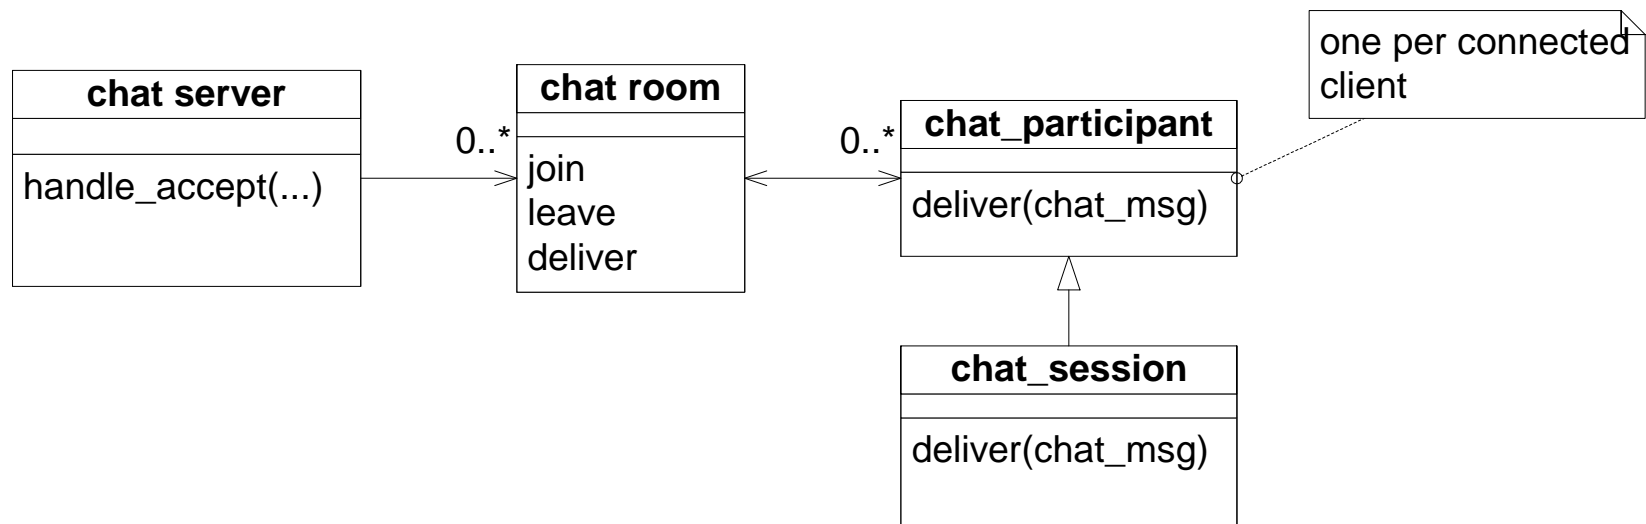
- ◆ Separates ‘listening behavior’ from client behavior
- ◆ New ‘ServiceHandler’ for each client
- ◆ Acceptor class is ‘factory’ for service handler
- ◆ May also handle any application specific setup protocol

# Chat Example – Network View

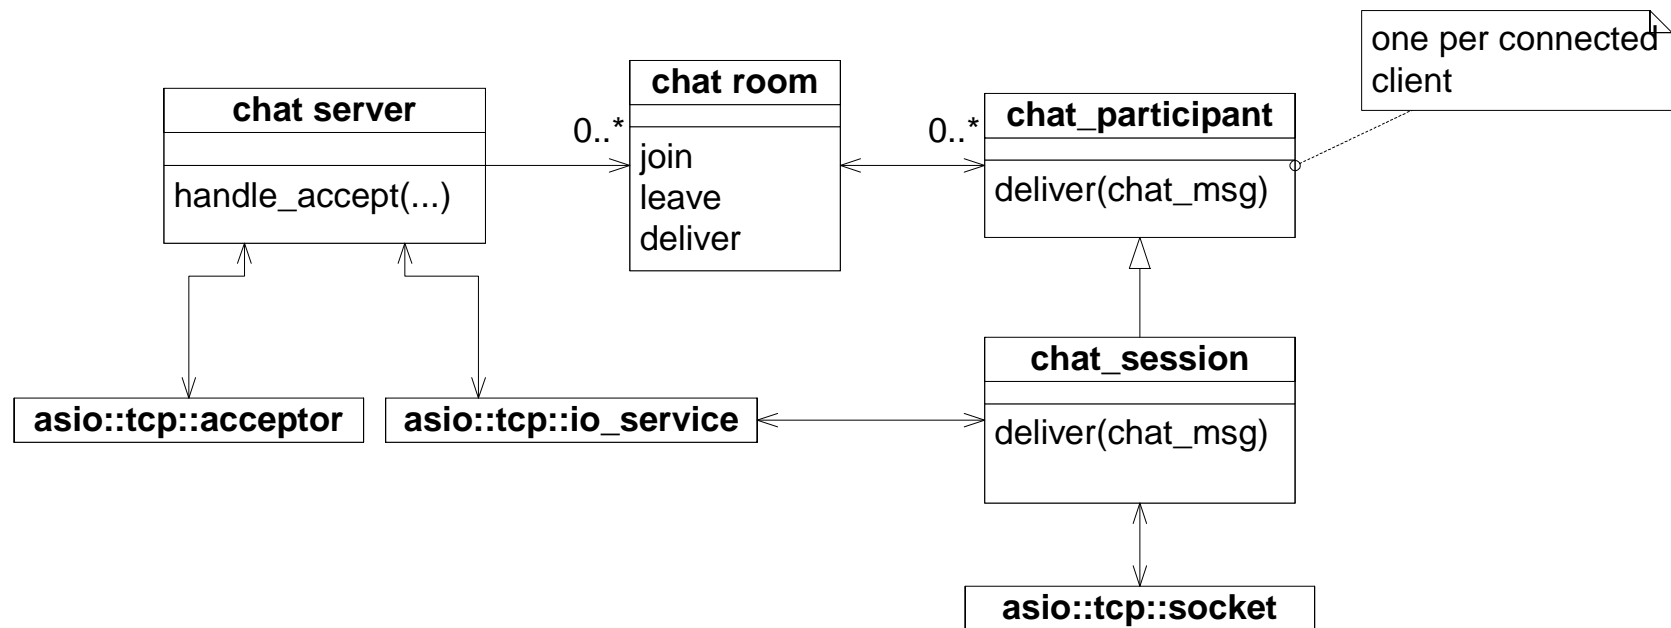




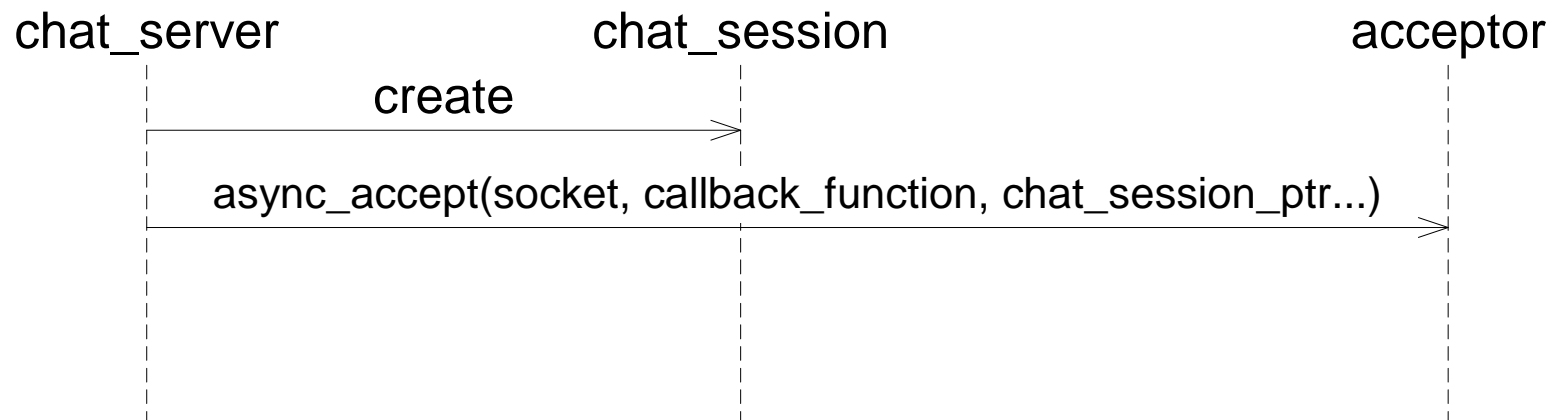
# Chat Classes Overview



# Chat Classes Detail



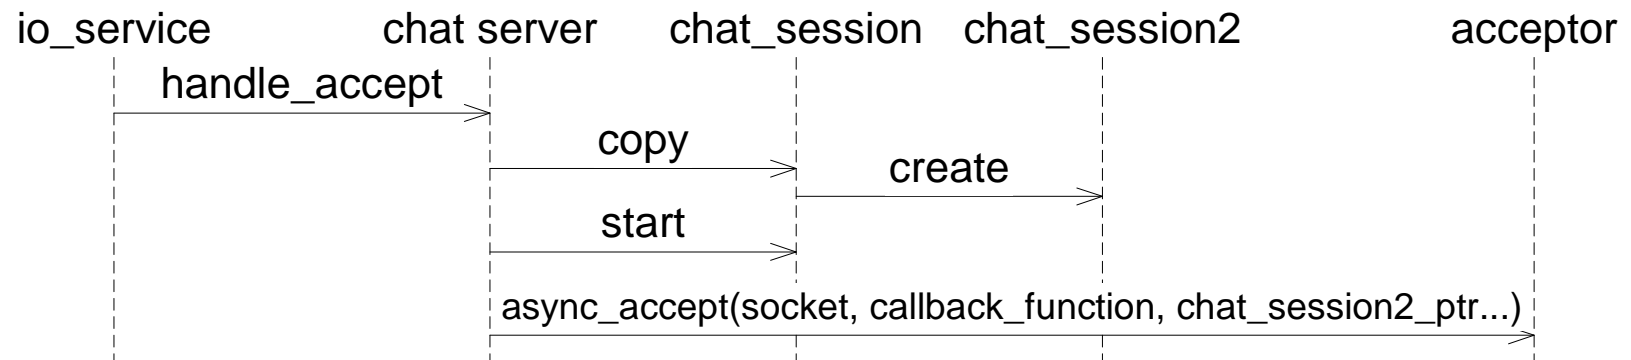
# Listening for Client Connection



# Listening for Client Connection

```
class chat_server
{
public:
    chat_server(boost::asio::io_service& io_service,
                const tcp::endpoint& endpoint)
        : io_service_(io_service),
          acceptor_(io_service, endpoint)
    {
        chat_session_ptr new_session(new chat_session(io_service_, room_));
        acceptor_.async_accept(new_session->socket(),
                               boost::bind(&chat_server::handle_accept, this, new_session,
   boost::asio::placeholders::error));
    }
};
```

# Session Setup Sequence




# Session Setup Code

```
class chat_server
{

void handle_accept(chat_session_ptr session,
                   const boost::system::error_code& error)
{
    if (!error)
    {
        session->start();
        chat_session_ptr new_session(new chat_session(io_service_, room_));
        acceptor_.async_accept(new_session->socket(),
                               boost::bind(&chat_server::handle_accept, this, new_session,
                               boost::asio::placeholders::error));
    }
}
```

# Chat - Current State

- ◆ Single established socket connection



# Chat Session Start

- ◆ Wait for inbound socket data
- ◆ Setup async read for header



# Chat – Handling Inbound Message – header

```
void handle_read_header(const boost::system::error_code& error)
{
    if (!error && read_msg_.decode_header())
    {
        boost::asio::async_read(socket_,
            boost::asio::buffer(read_msg_.body(), read_msg_.body_length()),
            boost::bind(&chat_session::handle_read_body, shared_from_this(),
                boost::asio::placeholders::error));
    }
    else
    {
        room_.leave(shared_from_this());
    }
}
```

# Chat – Handling Inbound Message - Body

```
void handle_read_body(const boost::system::error_code& error)
{
    if (!error)
    {
        room_.deliver(read_msg_);
        boost::asio::async_read(socket_,
            boost::asio::buffer(read_msg_.data(), chat_message::header_length),
            boost::bind(&chat_session::handle_read_header, shared_from_this(),
                boost::asio::placeholders::error));
    }
    else
    {
        room_.leave(shared_from_this());
    }
}
```

# Chat – sending outbound

```
class chat_room {  
    //...  
    void deliver(const chat_message& msg)  
    {  
        //...  
        std::for_each(participants_.begin(), participants_.end(),  
                       boost::bind(&chat_participant::deliver, _1, boost::ref(msg)));  
    }  
}
```

# Chat – sending outbound

```
void deliver(const chat_message& msg)
{
    bool write_in_progress = !write_msgs_.empty();
    write_msgs_.push_back(msg);
    if (!write_in_progress)
    {
        boost::asio::async_write(socket_,
            boost::asio::buffer(write_msgs_.front().data(),
                write_msgs_.front().length()),
            boost::bind(&chat_session::handle_write, shared_from_this(),
                boost::asio::placeholders::error));
    }
}
```

# Chat – Sending Outbound

```
void handle_write(const boost::system::error_code& error)
{
    if (!error)
    {
        write_msgs_.pop_front();
        if (!write_msgs_.empty())
        {
            boost::asio::async_write(socket_,
                boost::asio::buffer(write_msgs_.front().data(),
                    write_msgs_.front().length()),
                boost::bind(&chat_session::handle_write, shared_from_this(),
                    boost::asio::placeholders::error));
        }
    }
    else
    {
        room_.leave(shared_from_this());
    }
}
```

# Chat Server – Main Program

```
//main...
try
{
    //...snip...
    boost::asio::io_service io_service;

    chat_server_list servers;
    for (int i = 1; i < argc; ++i)
    {
        using namespace std; // For atoi.
        tcp::endpoint endpoint(tcp::v4(), atoi(argv[i]));
        chat_server_ptr server(new chat_server(io_service, endpoint));
        servers.push_back(server);
    }

    io_service.run();
}
//catch, exit, etc
```

# Notes on Chat Concurrency

- ◆ Single threaded
- ◆ Will support many active clients
- ◆ Callback functions are ‘quick’
- ◆ Excellent example of how to build many basic servers

# Other ways to implement Chat

- ◆ Wouldn't need to use point to point communication
- ◆ Could use UDP multi-cast
- ◆ Might not need a 'server' at all – peer to peer



# Session Agenda

- ◆ Introduction
- ◆ Boost.asio Overview
- ◆ Basic Synchronous Client Program
- ◆ Boost.asio in depth
- ◆ Asynchronous Server Program
- ◆ Boost.serialization – short overview
- ◆ Pitfalls of Network Programming
- ◆ Network Programming Tools
- ◆ Resources and Conclusion

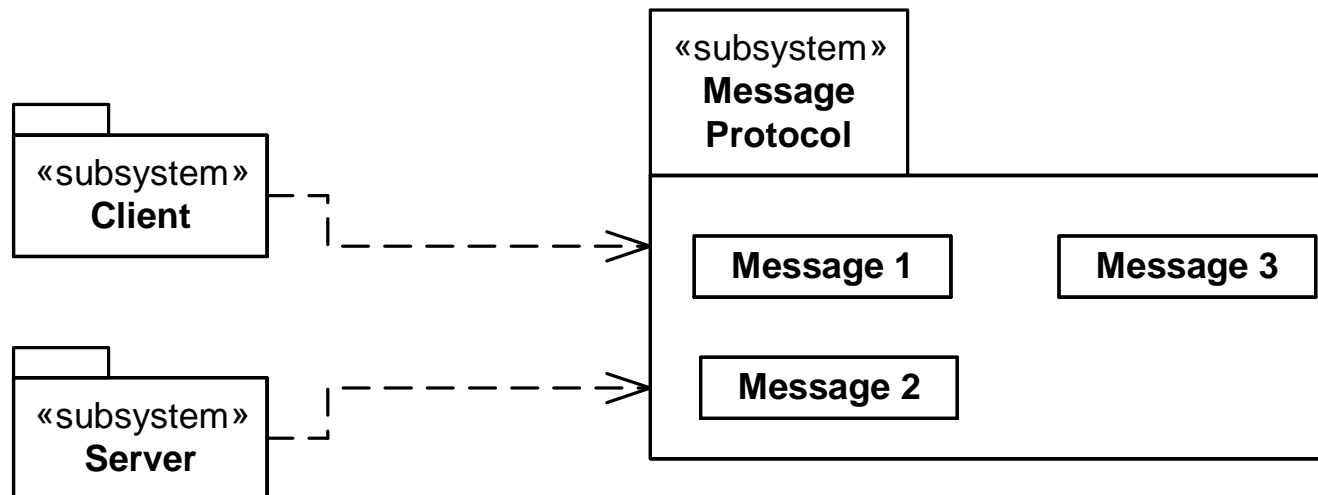

# Task: Send arbitrary Graph of Objects Between Processes

- ◆ Ultimately this is the task required in many network programs
- ◆ Implications
  - User defined types must be supported
  - Collections support
  - Must work with various platforms/compilers

# Serialization Design

- ◆ Serialization is Design pattern from long ago
- ◆ Separation of Concerns:
  - Object layout from format of ‘stream’
- ◆ Useful for network programs where data layout is a problem

# Messaging Architecture with Serialization



# serialization 101 – What is Boost.Serialization?

- ◆ Provides the ability to take data structures and make into a sequence of bytes
- ◆ Users can provide functions to serialize their own data types
- ◆ Automatically handles STL collections
- ◆ Many Boost libraries support
  - Date-time, multi-index, shared\_ptr, etc
- ◆ Author: Robert Ramey

# serialization 101 – history

- ◆ Initial version submitted in Feb 2002
- ◆ Initial review Nov 2002 – rejected
- ◆ Second review April 2004 – accepted
- ◆ Included in 1.32 release
- ◆ Details:  
<http://www.boost.org/libs/serialization/doc/index.html>

# serialization 101 – design goals

- ◆ Non-intrusive serialization
- ◆ Support for different archive formats (xml, binary)
- ◆ Archives can be extended by users
- ◆ Support for different stream types
- ◆ Versioning of classes

# serialization 101 – platforms

- ◆ Summary – modern compilers preferred
- ◆ VC6 (limits), VC7.1, VC8
- ◆ Intel 8+
- ◆ Borland (some limits)
- ◆ g++ 3.x + (full support – wchar issues on mingw)
- ◆ Sun 5.8
- ◆ HP-UX acc6 v2



# serialization 101 – Built Library

- ◆ Serialization is built into libraries
- ◆ To utilize will need to link appropriate libs

# Serializaton 101 - Archive Concept: Renders the data into an iostream

- ◆ Several different formats
  - Text based: 'text', 'xml'
  - Binary: 'binary'
- ◆ Archives can be 'loading' or 'saving'
- ◆ Archive class naming:
  - 'format'\_'chartype'\_'direction'archive
  - Where 'chartype' == blank (narrow) || 'w' wide
  - Where 'direction' == 'i' in || 'o' out
  - Xml, wide char, output ==  
boost::archive::xml\_woarchive

# serialization 101 – save a value

```
#include <boost/archive/text_oarchive.hpp>
#include <boost/date_time/gregorian/greg_serialize.hpp>

const date d(2007, May, 14);
std::ostringstream archive_stream;
boost::archive::text_oarchive archive(archive_stream);
archive << d; //2007-May-14 now in archive_stream buffer
```

# serialization 101 – serialization for custom types

```
#include <boost/date_time/gregorian/greg_serialize.hpp>
```

```
#include <boost/serialization/vector.hpp>
```

```
#include <boost/serialization/string.hpp>
```

```
class MyMessageType {  
    vector<int> vi;  
    date d;  
    string s;  
    friend class boost::serialization::access;  
    template<class Archive>  
    serialize(Archive& ar, const unsigned int version) const  
    {  
        ar & s;  
        ar & vi;  
        ar & d; //d has custom serialization method  
    }  
    public:  
        //constructors and useful functions  
};
```

# serialization 101 – serialization for custom types

```
template<class Archive>
serialize(Archive& ar, MyMessageType m &
          const unsigned int version) const
{
    ar & m.get_vi();
    ar & m.get_d();
    ar & m.get_s();
}
```

# serialization 101 – serialization for custom types - xml

```
template<class Archive>
serialize(Archive& ar, MyMessageType m &
          const unsigned int version) const
{
    ar & m.get_vi();
    ar & m.get_d();
    ar & m.get_s();
}
```

# Boost Serialization for Networked Programs

- ◆ Boost.serialization doesn't have a 'cross-platform' binary archive
- ◆ Need to stick to text-based archives such as xml or text archives
- ◆ Serialization could be extend to support true binary formats
  - CDR (from CORBA)
  - XDR (an IETF binary data standard)
    - <http://www.faqs.org/rfcs/rfc1014.html>
  - YAML SOC project

# Tradeoffs

- ◆ Advantages of Serialization Approach
  - Separates rendering from data types
  - Handles complex cases including pointers
  - Ensures client/server ‘automatically’ in sync
- ◆ Disadvantages
  - Not as fast as custom code



# Inbound Message Handling Design

- Problem:  
Connection supports multiple types of messages – how do you serialize and call the write handler?

| Stock Trade                                                      |
|------------------------------------------------------------------|
| string symbol<br>double price<br>time sale_time<br>long quantity |

| Stock Quote                                           |
|-------------------------------------------------------|
| string symbol<br>double ask_price<br>double bid_price |

# Inbound Message Design

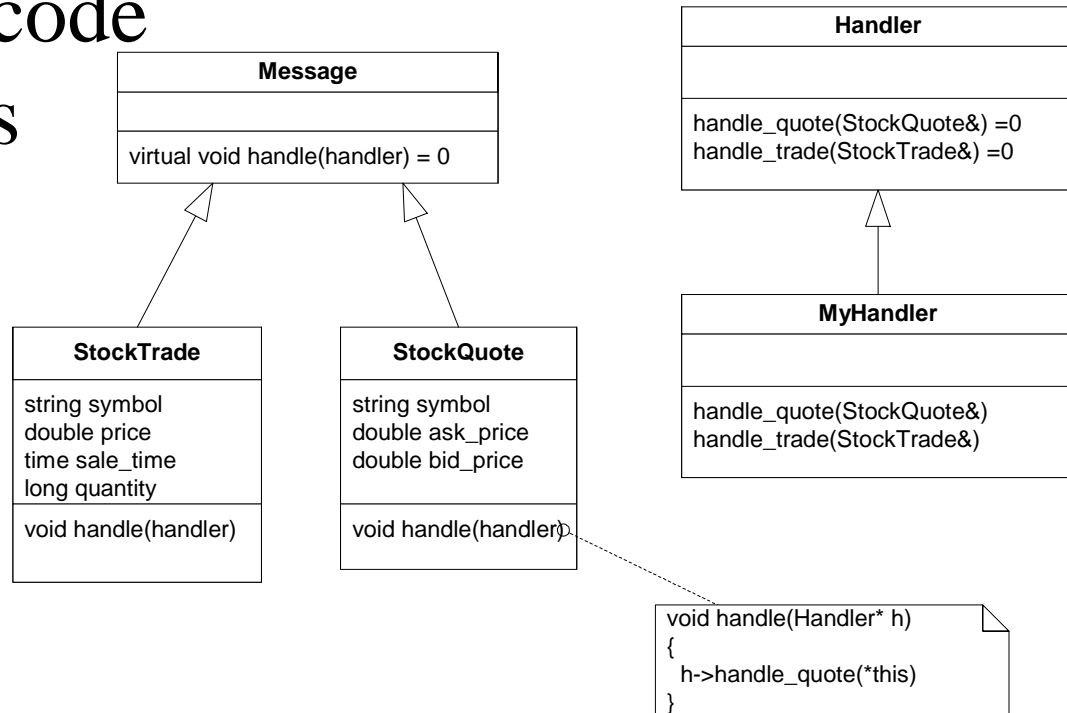
- ◆ Simple – add a message id to header then serialize correct type:

```
if (type == trade) {  
    StockTrade st;  
    ar >> st;  
    handle_trade(st);  
else if (type == quote) {  
    StockQuote sq;  
    ar >> sq;  
    handle_quote(sq);  
}
```

# Inbound Message Design

- Use polymorphism
- Serialization code never changes

```
Message m;  
ar >> m;  
m.handle();
```



# Session Agenda

- ◆ Introduction
- ◆ Boost.asio Overview
- ◆ Basic Synchronous Client Program
- ◆ Boost.asio in depth
- ◆ Asynchronous Server Program
- ◆ Boost.serialization – short overview
- ◆ Pitfalls of Network Programming
- ◆ Network Programming Tools
- ◆ Resources and Conclusion

# Pitfalls of Network Programming

- ◆ Assuming zero latency
  - But the speed of light is 186,000 mps?
- ◆ Assuming a reliable network
  - But TCP is reliable, right?
- ◆ Hard-coding addresses
  - But it's only a prototype?

# Pitfall 1: Latency in-depth

- ◆ Simple LAN
  - Latency is minimal ( < 0-10 milliseconds)
  - Most apps tested on local setup - dangerous
- ◆ WAN applications
  - Much larger latency – rule of thumb (< 10-100 milliseconds)
  - US/Europe theoretical min latency is ~ 30 milliseconds
- ◆ WAN over Satellite
  - Geostationary orbit - 22,300 miles (roundtrip 44, 600)
  - One hop == ~ .25 second delay – 250 milliseconds!
  - One round trip TCP sequence (send/ack) ~0.5 second delay
  - If one interaction requires 10 TCP request/responses that's 5 seconds!

# Latency Measurements – LAN

rtt min/avg/max/mdev = 38.058/41.519/51.352/4.112 ms

jeff@jeffdev2:~\$ **ping 192.168.1.100**

PING 192.168.1.100 (192.168.1.100) 56(84) bytes of data.

64 bytes from 192.168.1.100: icmp\_seq=1 ttl=128 **time=1.37 ms**

64 bytes from 192.168.1.100: icmp\_seq=2 ttl=128 **time=1.25 ms**

64 bytes from 192.168.1.100: icmp\_seq=3 ttl=128 **time=3.11 ms**

# Latency Measurements – Internet – ‘close’

```
jeff@jeffdev2:~$ ping -c 10 www.yahoo.com
```

```
PING www.yahoo-ht3.akadns.net (209.131.36.158) 56(84) bytes of data.
```

```
64 bytes from f1.www.vip.sp1.yahoo.com (209.131.36.158): icmp_seq=1 ttl=52 time=40.4 ms
```

```
64 bytes from f1.www.vip.sp1.yahoo.com (209.131.36.158): icmp_seq=2 ttl=52 time=42.9 ms
```

```
64 bytes from f1.www.vip.sp1.yahoo.com (209.131.36.158): icmp_seq=3 ttl=52 time=37.3 ms
```

```
64 bytes from f1.www.vip.sp1.yahoo.com (209.131.36.158): icmp_seq=4 ttl=52 time=50.6 ms
```

```
64 bytes from f1.www.vip.sp1.yahoo.com (209.131.36.158): icmp_seq=5 ttl=52 time=64.7 ms
```

```
64 bytes from f1.www.vip.sp1.yahoo.com (209.131.36.158): icmp_seq=6 ttl=52 time=44.4 ms
```

```
64 bytes from f1.www.vip.sp1.yahoo.com (209.131.36.158): icmp_seq=7 ttl=52 time=37.6 ms
```

```
64 bytes from f1.www.vip.sp1.yahoo.com (209.131.36.158): icmp_seq=8 ttl=52 time=40.0 ms
```

```
64 bytes from f1.www.vip.sp1.yahoo.com (209.131.36.158): icmp_seq=9 ttl=52 time=51.5 ms
```

```
64 bytes from f1.www.vip.sp1.yahoo.com (209.131.36.158): icmp_seq=10 ttl=52 time=40.4 ms
```

```
--- www.yahoo-ht3.akadns.net ping statistics ---
```

```
10 packets transmitted, 10 received, 0% packet loss, time 9055ms
```

```
rtt min/avg/max/mdev = 37.357/45.025/64.710/8.030 ms
```



# Latency Measurements – Internet – Arizona to Australia

```
jeff@jeffdev2:~$ ping -c 10 au.yahoo.com
```

```
PING p1.www.vip.aue.yahoo.com (203.84.217.32) 56(84) bytes of data.
```

```
64 bytes from p1.www.vip.aue.yahoo.com (203.84.217.32): icmp_seq=1 ttl=236 time=260 ms
```

```
64 bytes from p1.www.vip.aue.yahoo.com (203.84.217.32): icmp_seq=2 ttl=236 time=197 ms
```

```
64 bytes from p1.www.vip.aue.yahoo.com (203.84.217.32): icmp_seq=3 ttl=235 time=230 ms
```

```
64 bytes from p1.www.vip.aue.yahoo.com (203.84.217.32): icmp_seq=4 ttl=236 time=200 ms
```

```
64 bytes from p1.www.vip.aue.yahoo.com (203.84.217.32): icmp_seq=5 ttl=236 time=280 ms
```

```
64 bytes from p1.www.vip.aue.yahoo.com (203.84.217.32): icmp_seq=6 ttl=235 time=284 ms
```

```
64 bytes from p1.www.vip.aue.yahoo.com (203.84.217.32): icmp_seq=7 ttl=235 time=284 ms
```

```
64 bytes from p1.www.vip.aue.yahoo.com (203.84.217.32): icmp_seq=8 ttl=235 time=228 ms
```

```
64 bytes from p1.www.vip.aue.yahoo.com (203.84.217.32): icmp_seq=9 ttl=236 time=263 ms
```

```
64 bytes from p1.www.vip.aue.yahoo.com (203.84.217.32): icmp_seq=10 ttl=236 time=199 ms
```

```
--- p1.www.vip.aue.yahoo.com ping statistics ---
```

```
10 packets transmitted, 10 received, 0% packet loss, time 10549ms
```

```
rtt min/avg/max/mdev = 197.213/242.937/284.684/34.321 ms
```

# Latency Measurements – Internet – Arizona to UK

```
jeff@jeffdev2:~$ ping -c 10 uk.yahoo.com
```

```
PING www.euro.yahoo-eu1.akadns.net (217.146.186.51) 56(84) bytes of data.
```

```
64 bytes from www.vip.ird.yahoo.com (217.146.186.51): icmp_seq=1 ttl=49 time=185 ms
```

```
64 bytes from www.vip.ird.yahoo.com (217.146.186.51): icmp_seq=2 ttl=49 time=166 ms
```

```
64 bytes from www.vip.ird.yahoo.com (217.146.186.51): icmp_seq=3 ttl=49 time=174 ms
```

```
64 bytes from www.vip.ird.yahoo.com (217.146.186.51): icmp_seq=4 ttl=49 time=169 ms
```

```
64 bytes from www.vip.ird.yahoo.com (217.146.186.51): icmp_seq=5 ttl=49 time=163 ms
```

```
64 bytes from www.vip.ird.yahoo.com (217.146.186.51): icmp_seq=6 ttl=49 time=166 ms
```

```
64 bytes from www.vip.ird.yahoo.com (217.146.186.51): icmp_seq=7 ttl=49 time=164 ms
```

```
64 bytes from www.vip.ird.yahoo.com (217.146.186.51): icmp_seq=8 ttl=49 time=165 ms
```

```
64 bytes from www.vip.ird.yahoo.com (217.146.186.51): icmp_seq=9 ttl=49 time=167 ms
```

```
64 bytes from www.vip.ird.yahoo.com (217.146.186.51): icmp_seq=10 ttl=49 time=167 ms
```

```
--- www.euro.yahoo-eu1.akadns.net ping statistics ---
```

```
10 packets transmitted, 10 received, 0% packet loss, time 9042ms
```

```
rtt min/avg/max/mdev = 163.422/169.112/185.240/6.168 ms
```

# Designing for Latency

- ◆ New hardware isn't going to fix it – it's physics
- ◆ Consider using UDP (no acknowledgement)
- ◆ Minimize the number of application calls in protocol
- ◆ Example –
  - Client needs to retrieve 'user information'
  - Design 1: Remote User Object
    - Each user attribute (name, address) requires remote call
  - Design 2: Locally Cached User Object
    - One remote call, proxy makes remaining local

# Pitfall 2: Reliability in depth

- ◆ Nothing about network is really reliable
  - Theoretically impossible
    - Node crashes, router overloads, software bugs
  - tcp behaviors
    - Buffering
    - Window sizing (nagle algorithm) – cuts bandwidth automatically
- ◆ ‘Guaranteed message delivery’...not
  - Practically all schemes have limits
  - What if messages are ‘out of order’?
- ◆ Design for failure – it is inevitable!

# Designing for Failure

- ◆ Provide ‘graceful degradation mode’ if possible
  - Web browsers (cannot connect, offline mode)
- ◆ Work thru error cases
  - Connection based
    - Test connects when server not available
    - Disconnect processes at various times
    - Send periodic messages – ping messages
  - Messaging
    - Application level message sequencing
    - Servers may require retransmit facilities

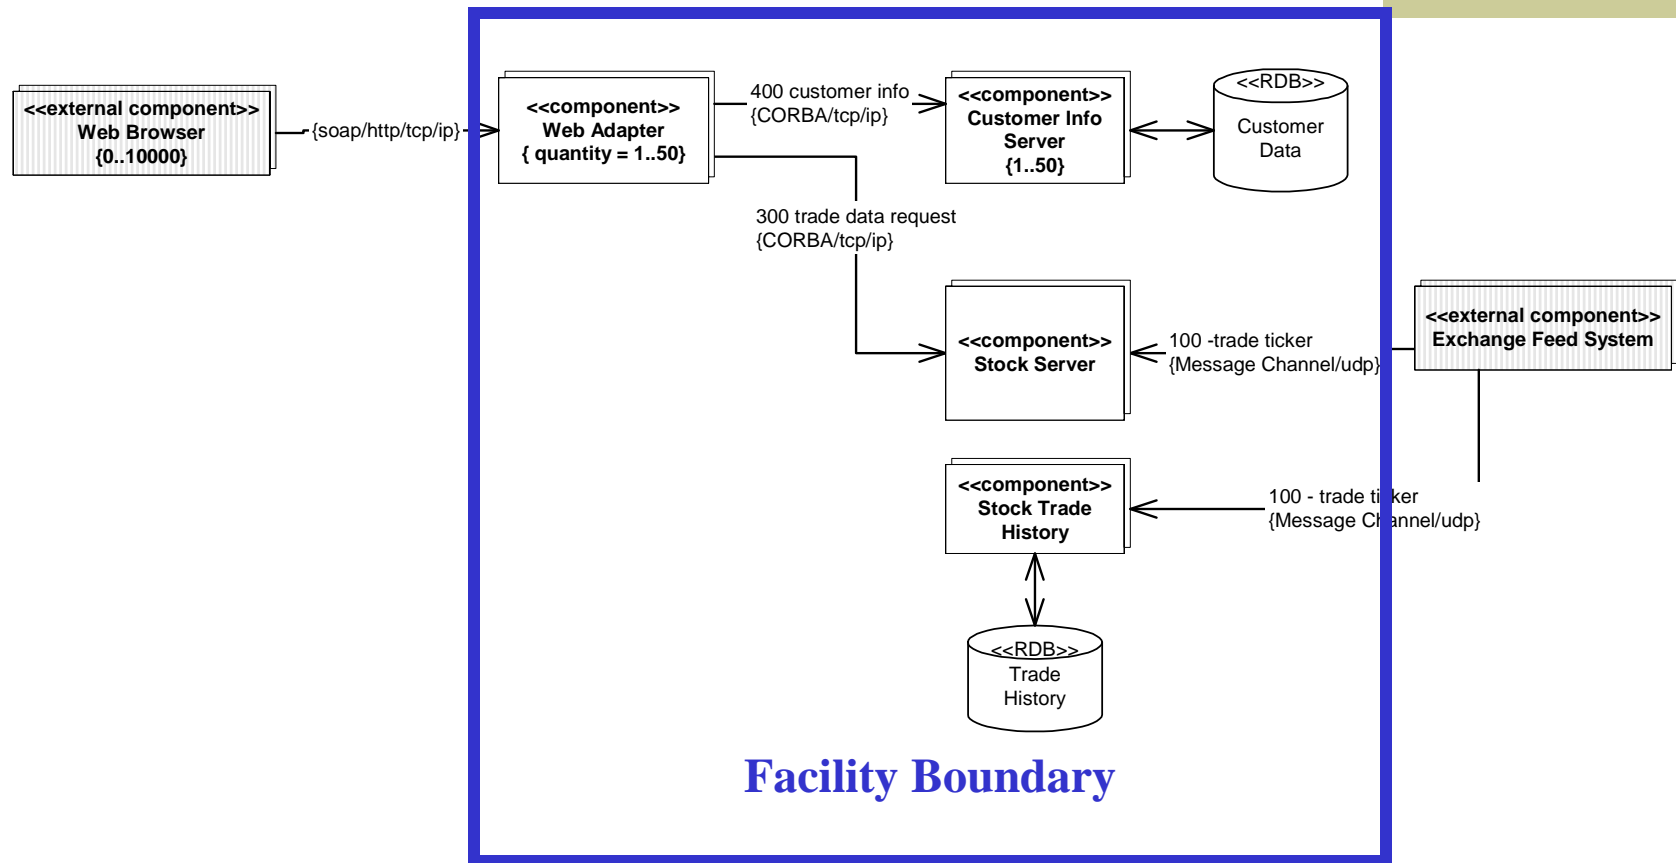
# Example: OMG DCPS QoS Specification

- ◆ DCPS – Data Centric Publish Subscribe
- ◆ Suitable for ‘real-time’ machines
- ◆ 23 different quality of service options
  - History (keep last, keep all) – kind of replay
  - Presentation – coherent data transmission
  - Reliability – kind of sequencing
  - etc

# Design for Failure – Startup Problem

- ◆ Design guideline
  - Order of process startup is completely arbitrary
  - Minimal amount of logic in each process
- ◆ Some functions may not be available until connectivity is complete
- ◆ Persistence often needed to all effective startup

# Startup Problem – An Example





# Session Agenda

- ◆ Introduction
- ◆ Boost.asio Overview
- ◆ Basic Synchronous Client Program
- ◆ Boost.asio in depth
- ◆ Asynchronous Server Program
- ◆ Boost.serialization – short overview
- ◆ Pitfalls of Network Programming
- ◆ Network Programming Tools
- ◆ Resources and Conclusion

# Tools: Traceroute Example

```
$ traceroute www.yahoo.com
```

Tracing route to www.yahoo-ht3.akadns.net [209.131.36.158]  
over a maximum of 30 hops:

```
 1  5 ms   9 ms   2 ms f1.www.vip.sp1.yahoo.com [209.131.36.158]
 2 29 ms  14 ms  14 ms 10.118.128.1
 3 24 ms  10 ms   7 ms ip68-2-6-110.ph.ph.cox.net [68.2.6.110]
 4 16 ms  15 ms  13 ms 68.2.13.94
 5 13 ms  17 ms  17 ms 68.2.13.9
 6 12 ms  32 ms  89 ms 68.2.13.5
 7 10 ms  13 ms  29 ms 68.2.13.1
 8 15 ms  13 ms  27 ms chnddsrj02-ae2.0.rd.ph.cox.net [68.2.14.5]
 9 32 ms  43 ms  32 ms langbbr01-ae0.r2.la.cox.net [68.1.0.232]
10 30 ms  25 ms  31 ms exchange-cust1.la1.equinix.net [206.223.123.16]
11 40 ms   *    38 ms ge-1-3-4-p142.pat1.pao.yahoo.com [216.115.96.42]
12 38 ms  42 ms  40 ms g-1-0-0-p150.msr2.sp1.yahoo.com [216.115.107.77]
13 41 ms  38 ms  42 ms UNKNOWN-209-131-32-23.yahoo.com [209.131.32.23]
14 45 ms  41 ms  43 ms f1.www.vip.sp1.yahoo.com [209.131.36.158]
```

Trace complete.

2008-May-09

Copyright© 2006-2008 CrystalClear Software  
Network Programming with Boost

# Tools: netstat (windows)

```
$ netstat --help
```

Displays protocol statistics and current TCP/IP network connections.

```
NETSTAT [-a] [-b] [-e] [-n] [-o] [-p proto] [-r] [-s] [-v] [interval]
```

- a Displays all connections and listening ports.
- b Displays the executable involved in creating each connection or listening port. In some cases well-known executables host multiple independent components, and in these cases the sequence of components involved in creating the connection or listening port is displayed. In this case the executable name is in [] at the bottom, on top is the component it called, and so forth until TCP/IP was reached. Note that this option can be time-consuming and will fail unless you have sufficient permissions.
- e Displays Ethernet statistics. This may be combined with the -s option.
- n Displays addresses and port numbers in numerical form.
- o Displays the owning process ID associated with each connection.
- v When used in conjunction with -b, will display sequence of components involved in creating the connection or listening port for all executables.

# Tools: netstat (unix)

usage: netstat [-veenNcCF] [<Af>] -r netstat {-V|--version|-h|--help}

netstat [-vnNcaeol] [<Socket> ...]

netstat { [-veenNac] -i | [-cnNe] -M | -s }

|                        |                                          |
|------------------------|------------------------------------------|
| -i, --interfaces       | display interface table                  |
| -g, --groups           | display multicast group memberships      |
| -v, --verbose          | be verbose                               |
| -n, --numeric          | don't resolve names                      |
| --numeric-hosts        | don't resolve host names                 |
| --numeric-ports        | don't resolve port names                 |
| --numeric-users        | don't resolve user names                 |
| -N, --symbolic         | resolve hardware names                   |
| -p, --programs         | display PID/Program name for sockets     |
| -l, --listening        | display listening server sockets         |
| -a, --all, --listening | display all sockets (default: connected) |
| -o, --timers           | display timers                           |

<Socket>={-t|--tcp} {-u|--udp} {-w|--raw} {-x|--unix} --ax25 --ipx --netrom

# Tools: netstat (windows)

```
$ netstat -b
```

## Active Connections

| Proto | Local Address                       | Foreign Address            | State       | PID   |
|-------|-------------------------------------|----------------------------|-------------|-------|
| TCP   | jeff_tablet:3314<br>[firefox.exe]   | localhost:3315             | ESTABLISHED | 9944  |
| TCP   | jeff_tablet:3315<br>[firefox.exe]   | localhost:3314             | ESTABLISHED | 9944  |
| TCP   | jeff_tablet:3322<br>[firefox.exe]   | localhost:3323             | ESTABLISHED | 9944  |
| TCP   | jeff_tablet:3323<br>[firefox.exe]   | localhost:3322             | ESTABLISHED | 9944  |
| TCP   | jeff_tablet:3450<br>[vncviewer.exe] | 192.168.1.102:5900         | ESTABLISHED | 11592 |
| TCP   | jeff_tablet:3487<br>[firefox.exe]   | cg-in-f167.google.com:http | ESTABLISHED | 9944  |

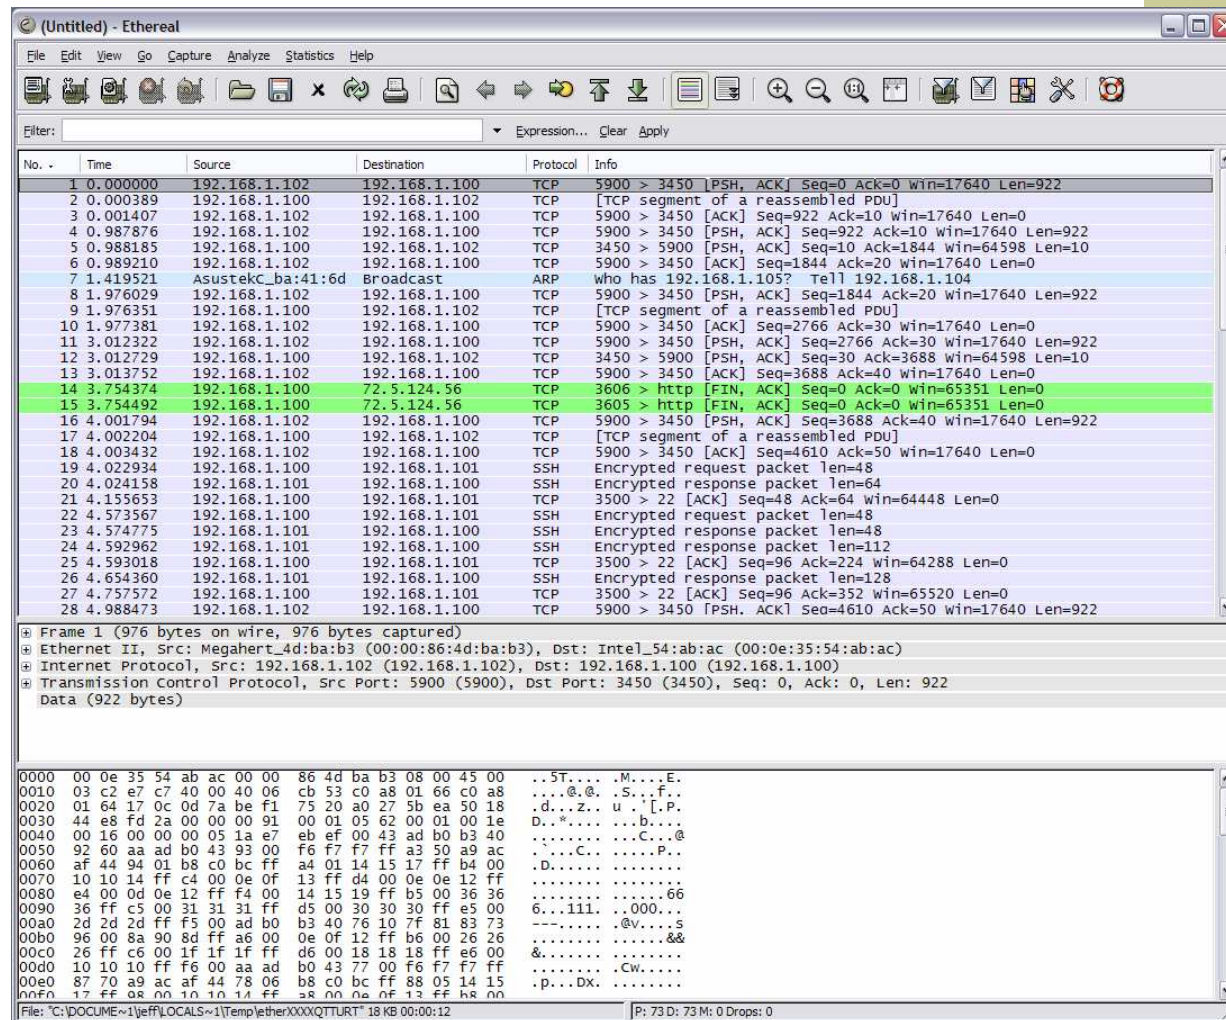
2008-May-09

Copyright© 2006-2008 CrystalClear Software  
Network Programming with Boost

# Ethereal/Wireshark

- ◆ Cross-platform network capture tool
  - Captures packets at os level
  - Great help in protocol debugging
  - Great learning tool
- ◆ Called ‘wireshark’ on some \*nix platforms
- ◆ <http://www.ethereal.com/>

# Tools: ethereal (windows)



2008-May-09

Copyright© 2006-2008 CrystalClear Software  
Network Programming with Boost

# Tools: ping

```
jeff@jeffdev2:~$ ping -c 10 www.yahoo.com
```

```
PING www.yahoo-ht3.akadns.net (209.131.36.158) 56(84) bytes of data.
```

```
64 bytes from f1.www.vip.sp1.yahoo.com (209.131.36.158): icmp_seq=1 ttl=52 time=38.3 ms
```

```
64 bytes from f1.www.vip.sp1.yahoo.com (209.131.36.158): icmp_seq=2 ttl=52 time=39.7 ms
```

```
64 bytes from f1.www.vip.sp1.yahoo.com (209.131.36.158): icmp_seq=3 ttl=52 time=40.6 ms
```

```
64 bytes from f1.www.vip.sp1.yahoo.com (209.131.36.158): icmp_seq=4 ttl=52 time=41.3 ms
```

```
64 bytes from f1.www.vip.sp1.yahoo.com (209.131.36.158): icmp_seq=5 ttl=52 time=47.2 ms
```

```
64 bytes from f1.www.vip.sp1.yahoo.com (209.131.36.158): icmp_seq=6 ttl=52 time=38.8 ms
```

```
64 bytes from f1.www.vip.sp1.yahoo.com (209.131.36.158): icmp_seq=7 ttl=52 time=38.0 ms
```

```
64 bytes from f1.www.vip.sp1.yahoo.com (209.131.36.158): icmp_seq=8 ttl=52 time=40.0 ms
```

```
64 bytes from f1.www.vip.sp1.yahoo.com (209.131.36.158): icmp_seq=9 ttl=52 time=51.3 ms
```

```
64 bytes from f1.www.vip.sp1.yahoo.com (209.131.36.158): icmp_seq=10 ttl=52 time=39.3 ms
```

```
--- www.yahoo-ht3.akadns.net ping statistics ---
```

```
10 packets transmitted, 10 received, 0% packet loss, time 9034ms
```

```
rtt min/avg/max/mdev = 38.058/41.519/51.352/4.112 ms
```



# Session Agenda

- ◆ Introduction
- ◆ Boost.asio Overview
- ◆ Basic Synchronous Client Program
- ◆ Boost.asio in depth
- ◆ Asynchronous Server Program
- ◆ Boost.serialization – short overview
- ◆ Pitfalls of Network Programming
- ◆ Network Programming Tools
- ◆ Resources and Conclusion

# Other Libraries of Interest

- ◆ Channel - Yigong Liu
  - Built on asio for Message Passing
  - <http://channel.sourceforge.net/>
- ◆ Signal Network (2007 SoC) – Stjepan Rajko
  - Has ‘remote signal’ concept
  - <http://dancinghacker.com/code/signet/index.html>
- ◆ Marshal/RPC - Stjepan Rajko
  - Factory of RPC/Marshalling part from above
  - <http://dancinghacker.com/code/marshal/index.html>

# Other Libraries of Interest

- ◆ RCF (Remote Call Framework) - **Jarl Lindrud**
  - Built on asio for C++ to C++ remote calls
  - [http://www.codeproject.com/threads/Rcf\\_Ipc\\_For\\_Cpp.asp](http://www.codeproject.com/threads/Rcf_Ipc_For_Cpp.asp)

# Other Libraries of Interest

## ◆ ACE/TAO

- Mature, portable C++ library includes threading, networking, shared memory, and CORBA implementations
- ACE == Adaptive Communications Environment
- TAO == Open source Object Request Broker
- <http://www.cs.wustl.edu/~schmidt/ACE.html>

# Boost.MPI – Message Passing Interface


- ◆ Message Passing Interface
  - C++ binding to MPI framework
  - Requires external MPI library
  - MPI supports many languages
  - MPI is traditionally used for ‘cluster computing’
- ◆ In Boost 1.35
- ◆ Uses boost.serialization for user defined types
- ◆ Authors: Douglas Gregor, Matthias Troyer

# Latency – Some Resources

- ◆ Read Waldo et. al.  
[http://research.sun.com/techrep/1994/sml\\_i\\_tr-94-29.pdf](http://research.sun.com/techrep/1994/sml_i_tr-94-29.pdf)
- ◆ AJAX Latency  
<http://richui.blogspot.com/2005/09/ajax-latency-problems-myth-or-reality.html>


# Network Programming Resources

- ◆ Unix Network Programming
  - <http://www.unpbook.com>
- ◆ Other Stevens books
  - TCP Illustrated
  - 3 volumes



# Schmidt C++ Books


- ◆ 2 volumes
- ◆ Describes details of using ACE
- ◆ Recommended if you're using ACE





# Schmidt Patterns

- ◆ Vol2 of Pattern Oriented Software Architectures
- ◆ Application design level networking descriptions



# Summary

- ◆ Network programming is ubiquitous
- ◆ New Boost libraries are making it easier than ever
- ◆ Go do it!