

Boost Date-Time

Authors Corner for BoostCon 2008

Jeff Garland



...Philosophy Moment...

Every generalization is dangerous,
especially this one.

-- Mark Twain



History and Stats

- Motivation: 1 project – too many time libraries
 - Differing needs – space efficiency versus resolution
 - 98% overlap
 - Convinced ‘OO’ design couldn’t work
- Stumbled across Boost around that time
- Reviewed and accepted into 1.29
 - Lacking a few semi-important features
- Later releases brought
 - format-based i/o facets
 - Time zones and `local_time`
- Library size
 - ~ 5000 statements (counted on semi-colons)
 - Test/Example code: ~6000 statements



Where is Date-Time Used Now?

- Many, many users worldwide
 - Who's Using Boost only scratches the surface
- Boost.Thread 1.35 uses duration types
- Boost.Asio uses duration types for timing
- Proposals to C++ Committee
 - TR2 for full date-time interfaces
 - C++0x threading interfaces



Library Design Requirements

- ❑ Basic Interfaces
 - Date, Time, Durations, Periods
 - Represented as valuetype – should behave like ‘integer’
- ❑ Calculation
 - Efficient and type safe
 - No hidden calculation error -- round off effects
- ❑ Representation
 - Different calendars
 - Different date epoch's an time precisions
- ❑ Clock Interfaces
- ❑ Various I/O Interfaces

Some Primary Concepts in Date-Time

- Point in Time

- 2008-May-04 00:00:00.0

- Duration

- 10 minutes
- 2 days

- Period

- [2008-May-04 00:00:00.0/2008-May-05 00:00:00.0)

- Clock

- Hardware device to get the current time from the computer

- Resolution


- Smallest representable time value

- Time Zones

- Assist in representing local times

- Calendar

- Maps time to 'label'
- Core calculation engine for representing years





Realization of Concepts

- Duration Types:
 - hours, minutes, seconds
 - microseconds, milliseconds, nanoseconds
 - **years, months**, weeks, days,
- Time Point Types:
 - date, ptime, local_time
- Period Types
 - date_period, time_period, local_time_period
- Iterator Type
 - date_iterator, time_iterator



From Concepts to Types

- Durations as types
 - Want to be able to represent durations in function interfaces and as data members
 - 1 set of durations types provide core library
- Time point types
- Durations and Time Points define the mathematical rules



Doing the Math - Examples

```
using namespace boost::gregorian;  
date weekstart(2002, Feb, 1);  
date weekend = weekstart + week(1);  
date d2 = d1 + days(5);  
  
date today = day_clock::local_day();  
if (d2 >= today) { } //date comparison operators
```



Doing the Math – Time Examples

```
using namespace boost::posix_time;
date d(2002, Feb, 1); //an arbitrary date
ptime t1( d, hours(5) + nanoseconds(100) );
//date + time of day offset
ptime t2 = t1 - minutes(4) + seconds(2);
ptime now = second_clock::local_time();
now += nanoseconds(100);
now -= days(1);
```



Doing the Math: Time Points

Defined:

Timepoint + Duration --> Timepoint

Timepoint - Duration --> Timepoint

Timepoint - Timepoint --> Duration

Undefined:

Duration + Timepoint --> Undefined

Duration - Timepoint --> Undefined

Timepoint + Timepoint --> Undefined



Doing the Math: Special Values

- ❑ Special Values – ‘not a date time’, infinities
- ❑ What happens when the user does arithmetic with these?

$\text{Timepoint}(\text{NADT}) + \text{Duration} \rightarrow \text{Timepoint}(\text{NADT})$

$\text{Timepoint}(\infty) + \text{Duration} \rightarrow \text{Timepoint}(\infty)$

$\text{Timepoint} + \text{Duration}(\infty) \rightarrow \text{Timepoint}(\infty)$

$\text{Timepoint} - \text{Duration}(\infty) \rightarrow \text{Timepoint}(-\infty)$



Doing the Math – Non Reversibility

- ❑ Unfortunately some time durations types aren't fixed size
- ❑ Leads to an issue with math reversibility
- ❑ Initially left these types out – added due to user request

```
date d(2005, Nov, 29);  
d += months(1); // "2005-Dec-29"  
d += months(1); // "2006-Jan-29"  
d += months(1); // "2006-Feb-28" --> snap-to-end-of-month behavior kicks in  
d += months(1); // "2006-Mar-31" --> unexpected result  
d -= months(4); // "2005-Nov-30" --> unexpected result, not where we started
```



Format Based I/O

```
date d(2004, Feb, 29);
time_duration td(12,34,56,789);
stringstream ss;
ss << d << ' ' << td; //"2004-Feb-29 12:34:56.000789"

ptime pt(not_a_date_time);
cout << pt << endl; // "not-a-date-time"
ss >> pt;
cout << pt << endl; // "2004-Feb-29 12:34:56.000789"

ss.str("");
ss << pt << " EDT-05EDT,M4.1.0,M10.5.0";
local_date_time ldt(not_a_date_time);
ss >> ldt;
cout << ldt << endl; // "2004-Feb-29 12:34:56.000789 EDT"
```



Facet Types and Format Specifiers

- Facet Type breakout
 - Wide and Narrow stream versions
 - Input and Output facets for each time point type
- Facets Provide “complete control” over input and output
- Date Facets
 - date_facet / wdate_facet
 - date_input_facet / wdate_input_facet
- Time Facets
 - time_facet / wtime_facet
 - time_input_facet / wtime_input_facet
- Local Time Facets
 - local_time_facet* / wlocal_time_facet*
 - local_time_input_facet* / wlocal_time_input_facet*



Using Format Strings - Localization

- Like strftime
 - Extensions for items not supported by strftime
- %x %X date/time format from the imbued locale.

```
ptime pt(date d(2005,Oct,31), hours(20));  
time_facet* f = new time_facet("%x %X");  
locale loc = locale(locale("en_US"), f);  
cout.imbue(loc);  
cout << pt; // "10/31/2005 08:00:00 PM"
```




Using Format Strings – ISO format

`%Y%m%dT%H%M%S%F%q`

`// Oct 15, 2005 13:12:11 MST`

`ISO: "20051015T131211-0700"`

□ Simple way – call methods to set iso format

`date_facet f = ...`

`f->set_iso_format(); // "%Y%m%d"`

`f->set_iso_extended_format(); // "%Y-%m-%d"`



Changing the Format

```
local_time_facet* output_facet = new local_time_facet();
local_time_input_facet* input_facet = new local_time_input_facet();
ss.imbue(locale(locale::classic(), output_facet));
ss.imbue(locale(ss.getloc(), input_facet));
output_facet->format("%a %b %d, %H:%M %z");
```

```
ss.str("");
ss << ldt;
cout << ss.str() << endl; // "Sun Feb 29, 12:34 EDT"
```

```
output_facet->format(local_time_facet::iso_time_format_specifier);
ss.str("");
ss << ldt;
cout << ss.str() << endl; // "20040229T123456.000789-0500"
```

```
output_facet->format(local_time_facet::iso_time_format_extended_specifier);
ss.str("");
ss << ldt;
cout << ss.str() << endl; // "2004-02-29 12:34:56.000789-05:00"
```



Date-Time: Flight Challenge

- ❑ Challenge: Program to calculate local arrival times for a cross country flight (Phoenix to New York) that flies through the daylight savings transition
- ❑ Flight is 4 hours and 30 minutes
- ❑ DST changed on April 2, 2005



User Code for Flight Challenge

```
typedef boost::shared_ptr<time_zone_base> zone_ptr;

//setup some timezones for creating and adjusting times
tz_database tz_db;
tz_db.load_from_file("date_time_zonespec.csv");
zone_ptr nyc_tz = tz_db.time_zone_from_region("America/New_York");
zone_ptr phx_tz(new posix_time_zone("MST-07:00:00"));

//local departure time in phoenix is 11 pm on april 2 2005
// (ny changes to dst on apr 3 at 2 am)
local_date_time phx_departure(date(2005, Apr, 2), hours(23), phx_tz,
                               local_date_time::NOT_DATE_TIME_ON_ERROR);

time_duration flight_length = hours(4) + minutes(30);
local_date_time phx_arrival = phx_departure + flight_length;
local_date_time nyc_arrival = phx_arrival.local_time_in(nyc_tz);
```



Hidden Gem: constrained_value

- Problem: Need to range check myriad of values on various interfaces.
 - month: 1..12
 - day: 1..31
- Solution: constrained_value template

- Stolen from Pascal
var x: 1..10;

- A bit more capable
 - Allows for customization of
- Not really so hidden
 - Dr. Dobbs article
<http://www.ddj.com/cpp/184401886>
 - Boost library proposal floating around



constrained_value template in detail

```
template<class value_policies>
class constrained_value {
public:
    typedef typename value_policies::value_type value_type;
    constrained_value(value_type value)
    {
        //does range checking – calls value_policies:on_error if
        //value is out of range
        assign(value);
    }
    constrained_value& operator=(value_type v) ; //calls assign
    static value_type max() {return (value_policies::max)();};
    static value_type min () {return (value_policies::min)();};
    operator value_type() const {return value_;};
```



Realization

- 3 typedefs define the rules for ‘months’

greg_month.hpp

// Exception thrown if a greg_month is constructed with a value out of range

struct bad_month : public std::out_of_range

{

bad_month() : std::out_of_range(std::string("Month number is out of range 1..12")) {}

};

// Build a policy class for the greg_month_rep

typedef CV::simple_exception_policy<unsigned short, 1, 12, bad_month> greg_month_policies;

// A constrained range that implements the gregorian_month rules

typedef CV::constrained_value<greg_month_policies> greg_month_rep;



Hidden Gems: Templates and Text

- ❑ Problem: Need a string constant as part of a library.
- ❑ Conventional wisdom
 - Definition in .cpp file
 - Build strings into a library to link
- ❑ Templates can allow header only implementation without ‘multiply defined’ linking problem



Conventional Approach

```
//date_facet.hpp
//mythical class for this example
class date_facet {
private:
    static const char default_period_separator[4];
//....
};
//...
```

```
//date_facet.cpp
date_facet::default_period_separator[4] = “ * “;
```



Template Solution

```
template <class date_type,  
        class CharT,  
        class OutItrT = std::ostreambuf_iterator<CharT, std::char_traits<CharT> > >  
class date_facet : public std::locale::facet {  
public:  
    static const char_type default_period_separator[4];  
    //....  
};  
//...
```

```
template <class date_type, class CharT, class OutItrT>  
const typename date_facet<date_type, CharT, OutItrT>::char_type  
date_facet<date_type, CharT, OutItrT>::default_period_separator[4] = { ' ', '/', ' ' };
```



Things I've Learned

- Valuetype programming is hard and under appreciated
 - Valuetypes often present large complex interfaces
 - Must be extremely efficient
- Writing good C++ I/O is hard
 - Extremely hard to understand all factors to write a truly standard operator<<
 - No one does it 'correctly'
 - Implementing facets and manipulators is painful
- Naming is hard
 - The ptime mistake
 - Needed to avoid 'time' because of conflicts on some older platforms
 - Should be date_time, but used that in the namespace
 - Too obscure
- Documentation best practice
 - Examples for every function
 - Date-time uses table
 - Function signature, description, example



Future Directions

- Duration Types
 - Will be refactored to match new standardization proposals
 - Replaces inheritance
- Standard
 - C++ 0x - ?
 - TR2
- 2008 Boost Summer of Code Project
 - New calendar types – historical and astronomical
- Boost.Timer integration – one of these centuries